


Homework#4

Recitation Exercises:

Problem 4:

- a) Probability = number of ways to select one centroid from each cluster / # of ways of selecting K clusters
 $P = \text{factorial}(k) / k^k$



- b) Probability that a sample of size $2K$ contains at least one points from each cluster.

$$P = 2 * \text{factorial}(k) / k^k$$

For $K = 10$

$$P = 2 * \text{factorial}(10) / 10^{10} = 0.000728$$

$$\text{For } k = 100, p = 1.867 * 10^{-42}$$

$$\text{For } k = 1000, p = \text{near to zero } (0)$$

Problem 7:

Let first see the squared error function:

$$SSE = \sum_{i=1}^K \sum_{x \in C_i} \text{dist}(C_i, x)^2$$

dist = standard Euclidean distance between two objects in Euclidean space. One easy way to reduce SSE is to increase K , the number of clusters.

Option C seems good reason that less dense region is normally consisting of noise and outliers. Also more closer the points higher the density. So, the result will be reduced SSE.

Problem 11:

if the SSE for one variable is low for all clusters?

- Then we have a variable which is constant and that has little use in dividing the data into group on clusters.

Low for just one cluster?

- It is the best case possible for any cluster as this attribute helps in defining this cluster.

High for all clusters?

- Then our data is scattered far away and contain only noise.

High for just one cluster?

- Then it is an odd with the information provided by the attributes with low SSE that define the cluster.

This means that cluster defined by this attribute are different from those defined by other attributes but it does not help define the cluster.

How could you use the per variable SSE information to improve your clustering?

- The idea is to eliminate attributes that have poor distinguishing power between clusters. The attributes with high SSE are troublesome if they have a relatively high SSE with respect to other attributes since they introduce a lot of noise into the computation of the overall SSE.

Problem 16:

Table:

	p1	p2	p3	p4	p5
p1	1.00	0.10	0.41	0.55	0.35
p2	0.10	1.00	0.64	0.47	0.98
p3	0.41	0.64	1.00	0.44	0.85
p4	0.55	0.47	0.44	1.00	0.76
p5	0.35	0.98	0.85	0.76	1.00


Single Linkage

First, we will cluster p2 & p5 together since distance is the highest (similarity)

	p1	p2, p5	p3	p4
p1	1.00	0.35	0.41	0.55
p2, p5	0.35	1.00	0.85	0.76
p3	0.41	0.85	1.00	0.44
p4	0.55	0.76	0.44	1.00


Next, cluster (p2, p5) with p3

	P_1	P_2, P_5, P_3	P_4
P_1	1.00	0.41	0.55
P_2, P_5, P_3	0.41	1.00	0.76
P_4	0.55	0.76	1.00




cluster, $(P_2, P_5), P_3$ with P_4 in the next step.

	P_1	P_2, P_5, P_3, P_4
P_1	1.00	0.55
P_2, P_5, P_3, P_4	0.55	1.00



1 will be the last to follow.



single link dendrogram

⇒ Complete Link:

Initially, join P_2 & P_5 to form a cluster.


	P_1	P_2, P_5	P_3	P_4		
P_1	1.00	0.10	0.41	0.55		
P_2, P_5	0.10	1.00	0.64	0.44	2	5
P_3	0.41	0.64	1.00	0.44		(0.98)
P_4	0.55	0.44	0.44	1.00		

Next, cluster P_2, P_5 with P_3 .

	P_1	P_2, P_5, P_3	P_4		
P_1	1.00	0.10	0.55		
P_2, P_5, P_3	0.10	1.00	0.44	2	5 3
P_4	0.55	0.44	1.00		

Cluster P_1 & P_4 together.

	P_1, P_4	P_2, P_5, P_3		
P_1, P_4	1.00	0.10	2	5 3 1 4
P_2, P_5, P_3	0.10	1.00		



Problem 17:

a) Data set: {6, 12, 18, 24, 30, 42, 48}

1) Centroids: {18, 45}

So, for 1st centroid 18 our cluster will be = {6, 12, 18, 24, 30}

$$SSE = (18-6)^2 + (18-12)^2 + (18-18)^2 + (18-24)^2 + (18-30)^2 = 360$$

For 2nd centroid 45 our cluster will be = {42, 48}

SSE = 18

Total SSE = 360 + 18 = **378**

2) Centroids: {15, 40}

1st SSE = 180

2nd SSE = 168

Total SSE = 180 + 168 = **348**

- b) I believe both centroid {18, 45} and {15, 40} represent stable solution. I do not think so we need any changes in the cluster generation.
- c) Two cluster produced by single link are {6, 12, 18, 24, 30} and {42, 48}
- d) MIN produces the most natural clustering in this solution that define cluster proximity as the proximity between the closest 2 points that are different clusters.
- e) This natural clustering corresponds to MIN that produce contiguous cluster. But center based is also taken as one of center given desired clusters.
- f) K mean is not good at finding clusters of different sizes at least when they are not well separated. The reason for this is that the objective of minimizing squared error causes it to break the large cluster.

Problem 21:

Entropy: entropy of a cluster $e_i = -\sum [(P_{ij}) \cdot \log_2(P_{ij})]$

Overall entropy $e = \sum [m_i/m \cdot e_i]$

Entropy_1 = $-[1/693 \cdot \log_2(1/693) + 1/693 \cdot \log_2(1/693) + 0 + 11/693 \cdot \log_2(11/693) + 4/693 \cdot \log_2(4/693) + 676/693 \cdot \log_2(676/693)]$

Entropy_1 = 0.20

Entropy_2 = 1.84

Entropy_3 = 1.70

Total Entropy = **1.44**

Purity: purity of a cluster $P_i = \max(P_{ij})$

Overall purity = $\sum (m_i/m \cdot P_i)$

Purity_1 = $676/693 = 0.98$

Purity_2 = $827/1562 = 0.53$

Purity_3 = $465/949 = 0.49$

Total Purity = **0.61**

Problem 22:

- a) Yes, the set of points that are uniformly spaced but have random arrangements will have regions of lesser or greater density. While the other set of points which are uniformly distributed over the unit square will have uniform density.
- b) For $K = 10$ cluster the uniformly distributed points will have a smaller SSE, because the points will be equally separated from the mean.
- c) DBSCAN will merge all points in one uniform data set into one cluster or classify them all as noise that will depend on its threshold boundary issues may arrive for the points that are at the edge. DBSCAN can find clusters in random data as well as it doesn't have validation in density.

Problem 1

```
In [2]: import pandas as pd
import numpy as np
from sklearn import preprocessing
from sklearn.preprocessing import Imputer
from sklearn.cluster import AgglomerativeClustering
```

```
In [5]: # data import for UIC
dataSet = pd.read_csv("http://archive.ics.uci.edu/ml/machine-learning-database
s/auto-mpg/auto-mpg.data",
                    delim_whitespace = True, header = None,
                    names = ['mpg', 'cylinders', 'displacement', 'horsepower', 'w
eight', 'acceleration', 'model', 'origin', 'car_name'])
```

```
In [9]: dataSet.describe()
```

Out[9]:

	mpg	cylinders	displacement	weight	acceleration	model	
count	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000	398
mean	23.514573	5.454774	193.425879	2970.424623	15.568090	76.010050	1.5
std	7.815984	1.701004	104.269838	846.841774	2.757689	3.697627	0.8
min	9.000000	3.000000	68.000000	1613.000000	8.000000	70.000000	1.0
25%	17.500000	4.000000	104.250000	2223.750000	13.825000	73.000000	1.0
50%	23.000000	4.000000	148.500000	2803.500000	15.500000	76.000000	1.0
75%	29.000000	8.000000	262.000000	3608.000000	17.175000	79.000000	2.0
max	46.600000	8.000000	455.000000	5140.000000	24.800000	82.000000	3.0

```
In [12]: dataSet.dtypes
```

```
Out[12]: mpg          float64
cylinders          int64
displacement      float64
horsepower        object
weight            float64
acceleration      float64
model             int64
origin            int64
car_name          object
dtype: object
```

```
In [13]: # see all unique values in horsepower because it is a object type col
dataSet.horsepower.unique()
```

```
Out[13]: array(['130.0', '165.0', '150.0', '140.0', '198.0', '220.0', '215.0',
                '225.0', '190.0', '170.0', '160.0', '95.00', '97.00', '85.00',
                '88.00', '46.00', '87.00', '90.00', '113.0', '200.0', '210.0',
                '193.0', '?', '100.0', '105.0', '175.0', '153.0', '180.0', '110.0',
                '72.00', '86.00', '70.00', '76.00', '65.00', '69.00', '60.00',
                '80.00', '54.00', '208.0', '155.0', '112.0', '92.00', '145.0',
                '137.0', '158.0', '167.0', '94.00', '107.0', '230.0', '49.00',
                '75.00', '91.00', '122.0', '67.00', '83.00', '78.00', '52.00',
                '61.00', '93.00', '148.0', '129.0', '96.00', '71.00', '98.00',
                '115.0', '53.00', '81.00', '79.00', '120.0', '152.0', '102.0',
                '108.0', '68.00', '58.00', '149.0', '89.00', '63.00', '48.00',
                '66.00', '139.0', '103.0', '125.0', '133.0', '138.0', '135.0',
                '142.0', '77.00', '62.00', '132.0', '84.00', '64.00', '74.00',
                '116.0', '82.00'], dtype=object)
```

```
In [14]: # replace '?' with np.NaN
dataSet['horsepower'] = dataSet['horsepower'].replace('?', np.nan)
```

```
In [15]: imputer_horsepower = Imputer(missing_values='NaN', strategy='mean', axis=0)

dataSet['horsepower'] = imputer_horsepower.fit_transform(dataSet['horsepower'
].values.reshape(-1,1),dataSet['horsepower'])
```

```
In [21]: # create new df with continuous variables for the dataset.
contus_df = pd.DataFrame([dataSet.mpg, dataSet.displacement, dataSet.horsepowe
r,dataSet.weight,dataSet.acceleration]).transpose()
```

```
In [40]: # create a copy
dataSet_copy = pd.DataFrame(dataSet)
dataSet_copy = dataSet_copy.drop(['car_name'], axis=1)
```

```
In [42]: # clustering
clustering = AgglomerativeClustering(n_clusters=3, affinity='euclidean', memor
y=None, connectivity=None,
                                     compute_full_tree =5, linkage ='average',
pooling_func = 'deprecated')
cluster = clustering.fit(contus_df)
```



```
In [43]: #Cluster labels
         clustering.fit_predict(contus_df)
```

[illegible]

```
In [47]: dataSet_copy['labels'] = cluster.labels_
dataSet_copy.groupby('origin').mean()
```

Out[47]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	mode
origin							
1	20.083534	6.248996	245.901606	118.814769	3361.931727	15.033735	75.61
2	27.891429	4.157143	109.142857	81.241983	2423.300000	16.787143	75.81
3	30.450633	4.101266	102.708861	79.835443	2221.227848	16.172152	77.44

```
In [48]: dataSet_copy.groupby('origin').var()
```

Out[48]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	mc
origin							
1	40.997026	2.760332	9702.612255	1569.532304	631695.128385	7.568615	13
2	45.211230	0.250311	509.950311	410.659789	240142.328986	9.276209	12
3	37.088685	0.348588	535.465433	317.523856	102718.485881	3.821779	13

```
In [49]: dataSet_copy['Clusters'] = cluster.labels_
dataSet_copy.groupby('Clusters').mean()
```

Out[49]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	mpg
Clusters							
0	27.365414	4.443609	131.934211	84.300061	2459.511278	16.298120	76
1	13.889062	8.000000	358.093750	167.046875	4398.593750	13.025000	73
2	17.510294	7.014706	278.985294	124.470588	3624.838235	15.105882	75

```
In [50]: dataSet_copy.groupby('Clusters').var()
```

Out[50]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	mpg
Clusters							
0	41.976309	0.851525	2828.083391	369.143491	182632.099872	5.718298	
1	3.359085	0.000000	2138.213294	756.521577	74312.340278	3.591429	
2	8.829892	1.059482	2882.492318	713.088674	37775.809263	10.556980	

Yes, we have a relationship between clusters assignment and class labels. We can see the mean values in both cases are similar. And we noticed that these are 3 classes in both labels as Origin and 3 clusters are created. We have a clear relationship.

Problem 2

```
In [98]: from sklearn import preprocessing, datasets
from sklearn.datasets import load_boston
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

```
In [99]: boston = load_boston()
# scaling the data using preprocessing.scale
scaled_boston_data = pd.DataFrame(data=preprocessing.scale(boston.data), columns= boston.feature_names)
```

```
In [100]: cluster_range = [2, 3, 4, 5, 6]

for n_clusters in cluster_range:
    km = KMeans(n_clusters = n_clusters, init='k-means++')
    cluster_labels = km.fit_predict(scaled_boston_data)
    silhouette_avg = silhouette_score(scaled_boston_data, cluster_labels)
    print(" For cluster= ", n_clusters, "Avg Silhouette is: ", silhouette_avg)
```

```
For cluster= 2 Avg Silhouette is: 0.359977342374
For cluster= 3 Avg Silhouette is: 0.257259122893
For cluster= 4 Avg Silhouette is: 0.280818056241
For cluster= 5 Avg Silhouette is: 0.276863962183
For cluster= 6 Avg Silhouette is: 0.285929370948
```

```
In [101]: # finding the mean for the optimum cluster = 2, greater
km = KMeans(n_clusters=2, init='k-means++')
cluster_labels = km.fit_predict(scaled_boston_data)
silhouette_avg = silhouette_score(scaled_boston_data, cluster_labels)
print(" For cluster= 2", "Avg Silhouette is: ", silhouette_avg)
scaled_boston_data['CLUSTER'] = cluster_labels
scaled_boston_data.groupby('CLUSTER').mean()
```

```
For cluster= 2 Avg Silhouette is: 0.359977342374
```

Out[101]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CLUSTER								
0	-0.388039	0.262392	-0.620368	0.002912	-0.584675	0.243315	-0.435108	0.468342
1	0.721270	-0.487722	1.153113	-0.005412	1.086769	-0.452263	0.808760	-0.279146

```
In [102]: data_summary = scaled_boston_data.loc[scaled_boston_data['CLUSTER'] == 0]
data_summary.describe()
```

Out[102]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	329.000000	329.000000	329.000000	329.000000	329.000000	329.000000	329.000000
mean	-0.388039	0.262392	-0.620368	0.002912	-0.584675	0.243315	-0.435108
std	0.045200	1.159853	0.622363	1.006458	0.534114	0.944020	0.954542
min	-0.417713	-0.487722	-1.557842	-0.272599	-1.465882	-1.868631	-2.335437
25%	-0.412070	-0.487722	-1.045700	-0.272599	-1.016689	-0.445397	-1.257951
50%	-0.405146	-0.487722	-0.720322	-0.272599	-0.576134	0.043261	-0.429393
75%	-0.387200	0.585267	-0.375976	-0.272599	-0.144217	0.734220	0.463180
max	-0.111580	3.804234	2.117615	3.668398	0.797361	3.476688	1.117494

```
In [103]: data_summary = scaled_boston_data.loc[scaled_boston_data['CLUSTER'] == 1]
data_summary.describe()
```

Out[103]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	
count	177.000000	1.770000e+02	177.000000	177.000000	177.000000	177.000000	177.000000
mean	0.721270	-4.877224e-01	1.153113	-0.005412	1.086769	-0.452263	0.808710
std	1.437545	1.447384e-15	0.310655	0.993564	0.718716	0.947526	0.406401
min	-0.407622	-4.877224e-01	1.015999	-0.272599	-0.196047	-3.880249	-1.005000
25%	-0.077985	-4.877224e-01	1.015999	-0.272599	0.512296	-0.935480	0.690000
50%	0.342909	-4.877224e-01	1.015999	-0.272599	1.073787	-0.290109	0.953000
75%	0.973109	-4.877224e-01	1.231945	-0.272599	1.367490	0.135863	1.074000
max	9.941735	-4.877224e-01	2.422565	3.668398	2.732346	3.555044	1.117000

```
In [104]: # scaling the data using preprocessing.normalize
scaled_boston_data_1 = pd.DataFrame(data=preprocessing.normalize(boston.data),
columns= boston.feature_names)
```

```
In [105]: cluster_range = [2, 3, 4, 5, 6]

for n_clusters in cluster_range:
    km = KMeans(n_clusters = n_clusters, init='k-means++')
    cluster_labels = km.fit_predict(scaled_boston_data_1)
    silhouette_avg = silhouette_score(scaled_boston_data_1, cluster_labels)
    print(" For cluster= ", n_clusters, "Avg Silhouette is: ", silhouette_avg)
    print(" Cluster Centroids= ", km.cluster_centers_)
```

For cluster= 2 Avg Silhouette is: 0.625790153315
Cluster Centroids= [[1.55650380e-02 -2.08166817e-17 2.58355757e-02
8.56414632e-05
9.51202648e-04 8.46013220e-03 1.27125744e-01 2.92385753e-03
2.96296276e-02 8.93089272e-01 2.78042327e-02 3.82724656e-01
2.60228732e-02]
[6.48561493e-04 3.23663943e-02 1.58698312e-02 1.47169613e-04
1.00436783e-03 1.28019832e-02 1.17982345e-01 9.07279657e-03
8.80792040e-03 6.06836448e-01 3.56496023e-02 7.72519906e-01
2.02319473e-02]]

For cluster= 3 Avg Silhouette is: 0.575733193467
Cluster Centroids= [[5.33766803e-04 3.61874259e-02 1.41978422e-02
1.49858483e-04
9.91875118e-04 1.30845203e-02 1.13438117e-01 9.67327874e-03
8.89716294e-03 5.88243515e-01 3.62365701e-02 7.89409879e-01
1.96933254e-02]
[9.94657937e-03 1.47186113e-03 2.59539951e-02 1.03376878e-04
9.67664212e-04 8.89467526e-03 1.30805856e-01 3.41404251e-03
2.25771863e-02 8.28056708e-01 2.81504699e-02 5.30891636e-01
2.42049719e-02]
[2.21562709e-02 -1.38777878e-17 2.75851760e-02 6.23949590e-05
1.04851698e-03 9.38784631e-03 1.38957760e-01 3.09816617e-03
3.35517426e-02 9.80050092e-01 3.06413005e-02 9.05811705e-02
3.11692578e-02]]

For cluster= 4 Avg Silhouette is: 0.492676766732
Cluster Centroids= [[1.35513597e-02 -2.08166817e-17 2.52396898e-02
9.27640141e-05
9.17225925e-04 8.13488591e-03 1.23238155e-01 2.85866780e-03
2.83372148e-02 8.64014925e-01 2.68483071e-02 4.81089806e-01
2.42978201e-02]
[4.48643313e-04 3.18358518e-02 1.44134947e-02 1.72013654e-04
1.01896098e-03 1.35363885e-02 1.20366189e-01 9.54134476e-03
8.83152187e-03 5.55872861e-01 3.70949081e-02 8.14295505e-01
1.97399010e-02]
[2.16604966e-02 -1.38777878e-17 2.76393386e-02 6.40813092e-05
1.05405111e-03 9.44466151e-03 1.38893581e-01 3.12118859e-03
3.35417959e-02 9.81098106e-01 3.06978454e-02 8.49706911e-02
3.12446556e-02]
[1.01061804e-03 3.33272194e-02 1.85072911e-02 1.02176466e-04
9.77939291e-04 1.14719579e-02 1.13665149e-01 8.22424473e-03
8.76517758e-03 6.99132708e-01 3.30321194e-02 6.96863311e-01
2.11230548e-02]]

For cluster= 5 Avg Silhouette is: 0.481830347736
Cluster Centroids= [[5.17499536e-04 1.26579884e-02 1.55039476e-02
1.83002829e-04
1.05364986e-03 1.34640067e-02 1.37595825e-01 8.60926443e-03
8.99760599e-03 5.51984876e-01 3.74438281e-02 8.16481356e-01
2.12790745e-02]
[2.16604966e-02 -1.38777878e-17 2.76393386e-02 6.40813092e-05
1.05405111e-03 9.44466151e-03 1.38893581e-01 3.12118859e-03
3.35417959e-02 9.81098106e-01 3.06978454e-02 8.49706911e-02
3.12446556e-02]
[1.38289801e-02 -2.08166817e-17 2.50398447e-02 9.53171521e-05
9.08185679e-04 8.00103924e-03 1.21291414e-01 2.76321427e-03
2.88441434e-02 8.66149000e-01 2.65703088e-02 4.78314411e-01
2.38635394e-02]
[1.38339654e-03 9.40306870e-03 2.28354236e-02 1.35170950e-04

```

1.04623160e-03 1.11212287e-02 1.37447650e-01 6.02373023e-03
8.84232300e-03 7.17563118e-01 3.30830741e-02 6.76803054e-01
2.45032866e-02]
[ 1.86221976e-04 1.00735088e-01 9.21225002e-03 7.75858853e-05
8.57199069e-04 1.33093828e-02 5.54432266e-02 1.35110102e-02
8.39463041e-03 6.04876152e-01 3.47523706e-02 7.80471992e-01
1.37253884e-02]]
For cluster= 6 Avg Silhouette is: 0.493266259425
Cluster Centroids= [[ 1.80084755e-04 9.06705111e-02 8.78666946e-03
1.35525272e-19
8.24908036e-04 1.27117756e-02 4.90516916e-02 1.36392169e-02
8.33882033e-03 6.44462556e-01 3.34048030e-02 7.51958680e-01
1.37781803e-02]
[ 2.16604966e-02 -1.38777878e-17 2.76393386e-02 6.40813092e-05
1.05405111e-03 9.44466151e-03 1.38893581e-01 3.12118859e-03
3.35417959e-02 9.81098106e-01 3.06978454e-02 8.49706911e-02
3.12446556e-02]
[ 1.83430627e-04 4.84832093e-02 1.43677352e-02 1.95490333e-04
1.02272793e-03 1.43545955e-02 1.10161603e-01 1.05306871e-02
7.96166416e-03 4.89884326e-01 3.92300179e-02 8.56570597e-01
1.86722971e-02]
[ 1.38289801e-02 -2.08166817e-17 2.50398447e-02 9.53171521e-05
9.08185679e-04 8.00103924e-03 1.21291414e-01 2.76321427e-03
2.88441434e-02 8.66149000e-01 2.65703088e-02 4.78314411e-01
2.38635394e-02]
[ 1.47320787e-03 1.49002359e-03 2.51285445e-02 1.34400862e-04
1.07852427e-03 1.09716164e-02 1.46093110e-01 5.10650984e-03
8.67461087e-03 7.23263447e-01 3.30444396e-02 6.70124146e-01
2.58147366e-02]
[ 7.14125095e-04 1.23567912e-02 1.50149355e-02 1.95591377e-04
1.04236448e-03 1.30515954e-02 1.41022577e-01 8.31673780e-03
9.63087845e-03 5.90149609e-01 3.62703464e-02 7.91211546e-01
2.13795741e-02]]

```

```

In [106]: # finding the mean for the optimum cluster = 2, greater
km = KMeans(n_clusters=2, init='k-means++')
cluster_labels = km.fit_predict(scaled_boston_data_1)
silhouette_avg = silhouette_score(scaled_boston_data_1, cluster_labels)
print(" For cluster= 2", "Avg Silhouette is: ", silhouette_avg)
scaled_boston_data_1['CLUSTER'] = cluster_labels
scaled_boston_data_1.groupby('CLUSTER').mean()

```

For cluster= 2 Avg Silhouette is: 0.625790153315

Out[106]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CLUSTER								
0	0.000649	0.032366	0.015870	0.000147	0.001004	0.012802	0.117982	0.00907
1	0.015565	0.000000	0.025836	0.000086	0.000951	0.008460	0.127126	0.00292

```
In [107]: data_summary = scaled_boston_data_1.loc[scaled_boston_data_1['CLUSTER'] == 0]
data_summary.describe()
```

Out[107]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	A
count	357.000000	357.000000	357.000000	357.000000	357.000000	357.000000	357.000000
mean	0.000649	0.032366	0.015870	0.000147	0.001004	0.012802	0.117982
std	0.001011	0.052953	0.011056	0.000527	0.000170	0.001825	0.054527
min	0.000013	0.000000	0.000963	0.000000	0.000660	0.008357	0.006437
25%	0.000119	0.000000	0.008207	0.000000	0.000888	0.011632	0.070957
50%	0.000243	0.000000	0.012721	0.000000	0.001003	0.012712	0.127134
75%	0.000663	0.044981	0.019558	0.000000	0.001070	0.013769	0.164100
max	0.007119	0.207644	0.061249	0.002282	0.001659	0.018156	0.237626

```
In [108]: data_summary = scaled_boston_data_1.loc[scaled_boston_data_1['CLUSTER'] == 1]
data_summary.describe()
```

Out[108]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	149.000000	149.0	149.000000	149.000000	149.000000	149.000000	149.000000	149.000000
mean	0.015565	0.0	0.025836	0.000086	0.000951	0.008460	0.127126	0.127126
std	0.017421	0.0	0.004999	0.000349	0.000224	0.001765	0.028194	0.028194
min	0.000129	0.0	0.020582	0.000000	0.000681	0.004682	0.051982	0.051982
25%	0.006232	0.0	0.023208	0.000000	0.000855	0.007464	0.116167	0.116167
50%	0.010680	0.0	0.023743	0.000000	0.000897	0.008157	0.125861	0.125861
75%	0.018505	0.0	0.026696	0.000000	0.000996	0.009070	0.131503	0.131503
max	0.113081	0.0	0.046424	0.002371	0.002065	0.015990	0.222496	0.222496

Silhouette score tells us the similarity an object is to its own cluster compared with another cluster. Here k=2, the silhouette score is the highest and thus optimal. The mean values for all features in each cluster for the optimal clustering is the same as that of the centroid co-ordinates.

Problem 3


```
In [125]: from sklearn.datasets import load_wine
from sklearn.cluster import KMeans
from sklearn.metrics.cluster import homogeneity_score, completeness_score
```

```
In [126]: data_wine = load_wine()
# scaling
scaled_wine_data = pd.DataFrame(data = preprocessing.scale(data_wine.data), columns=data_wine.feature_names)
```

```
In [127]: # for cluster = 3
cluster_model = KMeans(n_clusters=3, init = 'k-means++')
cluster_predict = cluster_model.fit_predict(scaled_wine_data)

target_values = data_wine.target
```

```
In [128]: #silhouette score
silhouette_avg = silhouette_score(scaled_wine_data, cluster_predict)
print(" For cluster= 3", "Avg Silhouette is: ", silhouette_avg)
```

For cluster= 3 Avg Silhouette is: 0.28485891919

```
In [129]: # cal Homogeneity/Completeness

h_score = homogeneity_score(target_values, cluster_predict)
print(h_score)
c_score = completeness_score(target_values, cluster_predict)
print(c_score)
```

0.878843200366

0.872963601608

```
In [130]: # normalize
norm_wine_data = pd.DataFrame(data = preprocessing.normalize(data_wine.data), columns=data_wine.feature_names)
```

```
In [131]: # for cluster = 3
cluster_model = KMeans(n_clusters=3, init = 'k-means++')
cluster_predict = cluster_model.fit_predict(norm_wine_data)

target_values = data_wine.target
```

```
In [132]: #silhouette score
silhouette_avg = silhouette_score(norm_wine_data, cluster_predict)
print(" For cluster= 3", "Avg Silhouette is: ", silhouette_avg)
```

For cluster= 3 Avg Silhouette is: 0.523346128229

```
In [133]: # cal Homogeneity/Completeness

h_score = homogeneity_score(target_values, cluster_predict)
print(h_score)
c_score = completeness_score(target_values, cluster_predict)
print(c_score)

0.376177457958
0.38858465671
```

if (Homogeneity == 1), then each cluster contains only members of a single class.

if (Completeness == 1), then all the members of a given class are assigned to the same cluster.

Homogeneity and completeness both score are used to measure the quality of the cluster.